

# logistic\_regression

April 13, 2020

## 1 linear regression

```
[1]: # imports
import pandas as pd
import matplotlib.pyplot as plt

# this allows plots to appear directly in the notebook
%matplotlib inline
```

### 1.1 Example: Advertising Data

```
[2]: # read data into a DataFrame
data = pd.read_csv('http://www-bcf.usc.edu/~gareth/ISL/Advertising.csv',
    ↪index_col=0)
data.head()
```

```
[2]:
```

	TV	Radio	Newspaper	Sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
3	17.2	45.9	69.3	9.3
4	151.5	41.3	58.5	18.5
5	180.8	10.8	58.4	12.9

What are the **features**? - TV: advertising dollars spent on TV for a single product in a given market (in thousands of dollars) - Radio: advertising dollars spent on Radio - Newspaper: advertising dollars spent on Newspaper

What is the **response**? - Sales: sales of a single product in a given market (in thousands of widgets)


```
[3]: # print the shape of the DataFrame
data.shape
```

```
[3]: (200, 4)
```

There are 200 **observations**, and thus 200 markets in the dataset.

```
[4]: # visualize the relationship between the features and the response using
      ↪ scatterplots
fig, axs = plt.subplots(1, 3, sharey=True)
data.plot(kind='scatter', x='TV', y='Sales', ax=axs[0], figsize=(16, 8))
data.plot(kind='scatter', x='Radio', y='Sales', ax=axs[1])
data.plot(kind='scatter', x='Newspaper', y='Sales', ax=axs[2])
```

```
[4]: <matplotlib.axes._subplots.AxesSubplot at 0xc1a3908>
```



Let's use **Statsmodels** to estimate the model coefficients for the advertising data:

```
[5]: # this is the standard import if you're using "formula notation" (similar to R)
import statsmodels.formula.api as smf

# create a fitted model in one line
lm = smf.ols(formula='Sales ~ TV', data=data).fit()

# print the coefficients
lm.params
```

```
[5]: Intercept    7.032594
TV              0.047537
dtype: float64
```

## 1.2 Interpreting Model Coefficients

## 1.3 Using the Model for Prediction

Let's say that there was a new market where the TV advertising spend was **\$50,000**. What would we predict for the Sales in that market?

$$y = \beta_0 + \beta_1 x$$

$$y = 7.032594 + 0.047537 \times 50$$

```
[6]: # manually calculate the prediction
     7.032594 + 0.047537*50
```

```
[6]: 9.409444
```

Thus, we would predict Sales of **9,409 widgets** in that market.

Of course, we can also use Statsmodels to make the prediction:

```
[7]: # you have to create a DataFrame since the Statsmodels formula interface
     ↪ expects it
     X_new = pd.DataFrame({'TV': [50]})
     X_new.head()
```

```
[7]:    TV
     0  50
```

```
[8]: # use the model to make predictions on a new value
     lm.predict(X_new)
```

```
[8]: array([ 9.40942557])
```

## 1.4 Plotting the Least Squares Line

Let's make predictions for the **smallest and largest observed values of x**, and then use the predicted values to plot the least squares line:

```
[9]: # create a DataFrame with the minimum and maximum values of TV
     X_new = pd.DataFrame({'TV': [data.TV.min(), data.TV.max()]})
     X_new.head()
```

```
[9]:    TV
     0  0.7
     1 296.4
```


```
[10]: # make predictions for those x values and store them
preds = lm.predict(X_new)
preds
```

```
[10]: array([ 7.0658692 , 21.12245377])
```

```
[11]: # first, plot the observed data
data.plot(kind='scatter', x='TV', y='Sales')

# then, plot the least squares line
plt.plot(X_new, preds, c='red', linewidth=2)
```

```
[11]: [<matplotlib.lines.Line2D at 0x14625128>]
```



## 1.5 Confidence in our Model

```
[12]: # print the confidence intervals for the model coefficients
lm.conf_int()
```

```
[12]:
```

	0	1
Intercept	6.129719	7.935468
TV	0.042231	0.052843